



# Rensselaer

why not change the world?®

# Measuring Repeatable Read with Read/Write Locking Performance

# Agenda

1. Motivation

2. Transaction Isolation

3. Database Implementation

4. Repeatable Read Implementation

5. Results and Findings

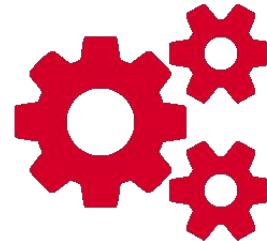


# Motivation



## Problem

- Data inconsistency in database systems with concurrent transactions
  - Multiple transactions trying to update the same data at the same time
  - One transaction trying to read data that has been updated by another transaction
- Need to ensure isolation between transactions using concurrency control techniques



## Solution

- Try to emulate the effect of running transactions in parallel using repeatable read
- Implement read/write locking to ensure data integrity



# Motivation

Database	
ID	Balance
0	10
1	10

- If run sequentially, Transaction 2 will always print \$20
- What happens if Transaction 1 finishes in the middle of Transaction 2?

Thread 1 - Transfers \$10 from ID 0 to 1

```
begin_tx  
a = readId(0, "Balance")  
b = readId(1, "Balance")  
writeId(0, a - 10)  
writeId(1, b + 10)  
commit_tx
```

Thread 2 - Adds Balance of ID 0 and 1

```
begin_tx  
a = readId(0, "Balance")  
b = readId(1, "Balance")  
print(a + b)  
commit_tx
```



# Motivation

Thread 1 - Transfers \$10 from ID 0 to 1

```
begin_tx  
a = readId(0, "Balance")  
b = readId(1, "Balance")  
writeId(0, a - 10)  
writeId(1, b + 10)  
commit_tx
```

Thread 2 - Adds Balance of ID 0 and 1

```
begin_tx  
a = readId(0, "Balance")  
b = readId(1, "Balance")  
print(a + b)  
commit_tx
```

	Thread 1	Thread 2
Timestep 1	commit_tx	a = readId(0, "Balance") → \$10
Timestep 2		b = readId(1, "Balance") → \$20
Timestep 3		print(a + b) → \$30 <b>Problem!</b>



# Transaction Isolation Levels

	<b>Isolation Level</b>	<b>Description</b>
Least Desirable	Read uncommitted	Allows seeing uncommitted data changes from other transactions
	Read Committed	Allows reading only committed data, but non-repeatable reads are possible
	Repeatable Read	Allows reading the same data repeatedly in the same transaction
	Serializable	Provides the highest level of isolation, guarantees correctness but with a performance penalty



# Transaction Design

## Custom transaction/query language

```
begin_tx
a = readId(0, "Balance")
b = readVal("Name", "George", "Balance")
if (a >= 10)
    writeId(0, a - 10)
    writeVal("Name", "George", b + 10)
else
    abort_tx
endif
commit_tx
```



# Transaction Design

## Custom transaction/query language

```
Begins the transaction      → begin_tx
Variable assignment        → a = readId(0, "Balance") ← Get column value based on ID
                            b = readVal("Name", "George", "Balance")
                            if (a >= 10)
                                writeId(0, a - 10)
                                writeVal("Name", "George", b + 10)
                            else
                                abort_tx
                            endif
                            commit_tx
```



# Transaction Design

## Custom transaction/query language

Begins the transaction → begin\_tx

Variable assignment → a = readId(0, "Balance") ← Get column value based on ID

If/else statements → if (a >= 10)  
writeId(0, a - 10)

Write to ID or Value → writeVal("Name", "George", b + 10) ← Arithmetic operations  
else  
abort\_tx  
endif  
commit\_tx



# Transaction Design

## Custom transaction/query language

```
Begins the transaction      begin_tx
Variable assignment        a = readId(0, "Balance")           Get column value based on ID
If/else statements          b = readVal("Name", "George", "Balance")
                           if (a >= 10)
                               writeId(0, a - 10)
                           else
                               writeVal("Name", "George", b + 10)   Arithmetic operations
Cancels transactions       abort_tx
Commit the transaction     commit_tx
```



# Database Implementation

---

- Current implementation stores data in void\* hash table of incrementing unique IDs
  - O(1) for data retrieval with O(n) for search
  - Large overhead to store data on disk compared to B-Tree
- Database defined based on custom schema
- Many robust database systems use B-Trees
  - Fast storage and retrieval of large amounts of data
  - Highly scalable and can easily adapt to increasing amounts of data
  - Minimize disk access by storing data in nodes that are optimized for disk blocks
  - Self-balancing as data is inserted



# Database Implementation

- Table design with unique IDs at every row
- Define table schemas with number of columns and their type

```
DatabaseTableSchema schema = {  
    .columnCount = 2,  
    .columnNames = new string[2] {"Name", "Balance"},  
    .columnTypes = new int[2] {DATABASE_TYPE_STRING, DATABASE_TYPE_INT}  
};
```

- Supported functions
  - `setID(ID, ColumnName, Value)` → Set a value in particular row, column
  - `setVal(SearchColumn, SearchValue, ColumnName, Value)` → Set a value in particular row, column by search
  - `getID(ID, ColumnName)` → Get a value in particular row, column
  - `getVal(SearchColumn, SearchValue, ColumnName)` → Get a value in particular row, column by search



# Repeatable Read

---

- Repeatable read: isolation level in database systems
  - Ensures transactions see the same data throughout the entire transaction
  - Other transactions modifying this data at the same time will not affect the transaction
- Used in applications that require a high degree of data consistency
- Improves performance by reducing number of conflicts between transactions
- Implemented with **locking**



## SHARED lock

- Transactions can read the same data simultaneously
- Prevents modifying the data
- Used when transaction only needs to *read* data

## EXCLUSIVE lock

- Prevents other transactions from reading or writing the locked data
- Used when a transaction needs to write/modify data
- Other transactions cannot lock this data



# Two-Phase Locking (2PL)

<b>Growing phase</b>	Locks are <i>acquired</i> , no locks are released
<b>Shrinking phase</b>	Locks are <i>released</i> , no locks are acquired

- Transaction must obtain all necessary locks before executing transactions
  - Must not release any locks until execution is completed
- Prevents conflicting transactions from overlapping in their execution
  - One transaction must wait for the other to finish before starting



# Importance of Locking

---

- **Data consistency**
  - Transactions will always see the same data, even if other transactions are updating the data
- **Deadlock prevention**
  - Occurs when two or more transactions are waiting for each other to release a lock
  - Locking ensures that transactions acquire locks in a specific order
- **Lost update prevention**
  - Occurs when a transaction updates data that has been updated by another transaction
  - Locking ensures transactions only update data that has not been updated by another transaction

# Transaction Manager: ReadWriteLockingTable

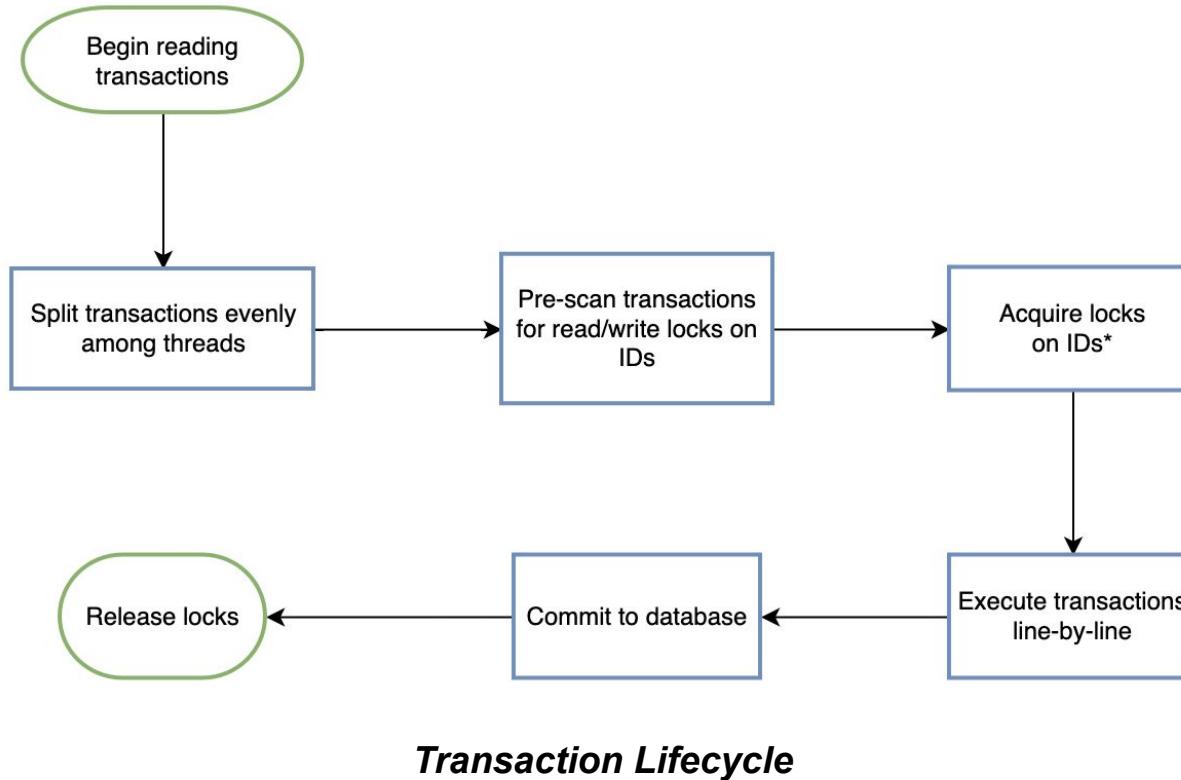
## ReadWriteLockingTable (unordered\_map)

Data ID	Read Locked?	Write Locked?	
0	No	Yes	 <i>SHARED lock</i>
57	No	No	 <i>Unlocked</i>
421	Yes	Yes	 <i>EXCLUSIVE lock</i>
...	...	...	

- **SHARED lock:** Write Locked = Yes
  - Used for read-only IDs
- **EXCLUSIVE lock:** Read Locked = Yes && Write Locked = Yes
  - Used for IDs that need to be modified



# Locking Data for Transactions



## Growing phase:

- Acquire all read/write locks before beginning execution

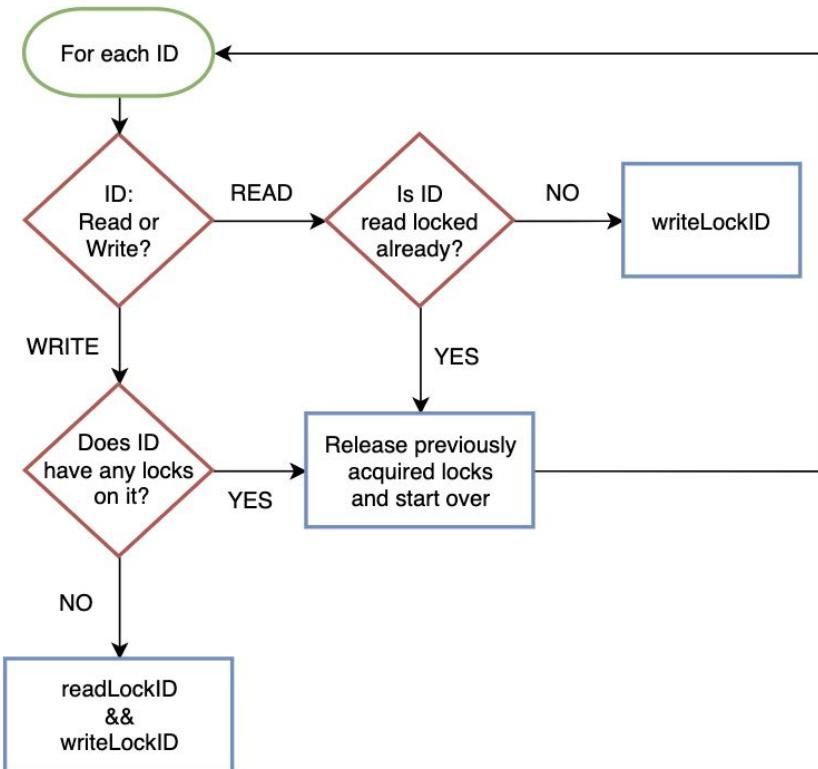
## Shrinking phase:

- Release all locks after transaction has been committed



# Acquiring Locks

- List of IDs that need to be locked generated by transaction pre-scan
- If conflict between transactions, wait for all necessary IDs to be unlocked
  - Function returns when all IDs have been successfully locked
- Avoids deadlocking
  - Tries to acquire all locks at once
  - Does not hold any locks indefinitely



## Repeatable Read: Other Implementations

Read/Write Locking	Snapshot Isolation	Optimistic Concurrency Control
<ul style="list-style-type: none"><li>• Acquire locks on data that transaction needs to access</li><li>• When transaction is complete, release locks</li><li>• High performance</li></ul>	<ul style="list-style-type: none"><li>• Transactions see data from a snapshot taken at the start of the transaction</li><li>• Modifications to database made after snapshot are not visible to transaction</li><li>• High storage and computational overhead</li></ul>	<ul style="list-style-type: none"><li>• Assume transactions will not conflict with each other</li><li>• Uses timestamps on data to check for conflicts</li><li>• If conflict, roll back transaction</li><li>• Used in environments with low data contention</li></ul>



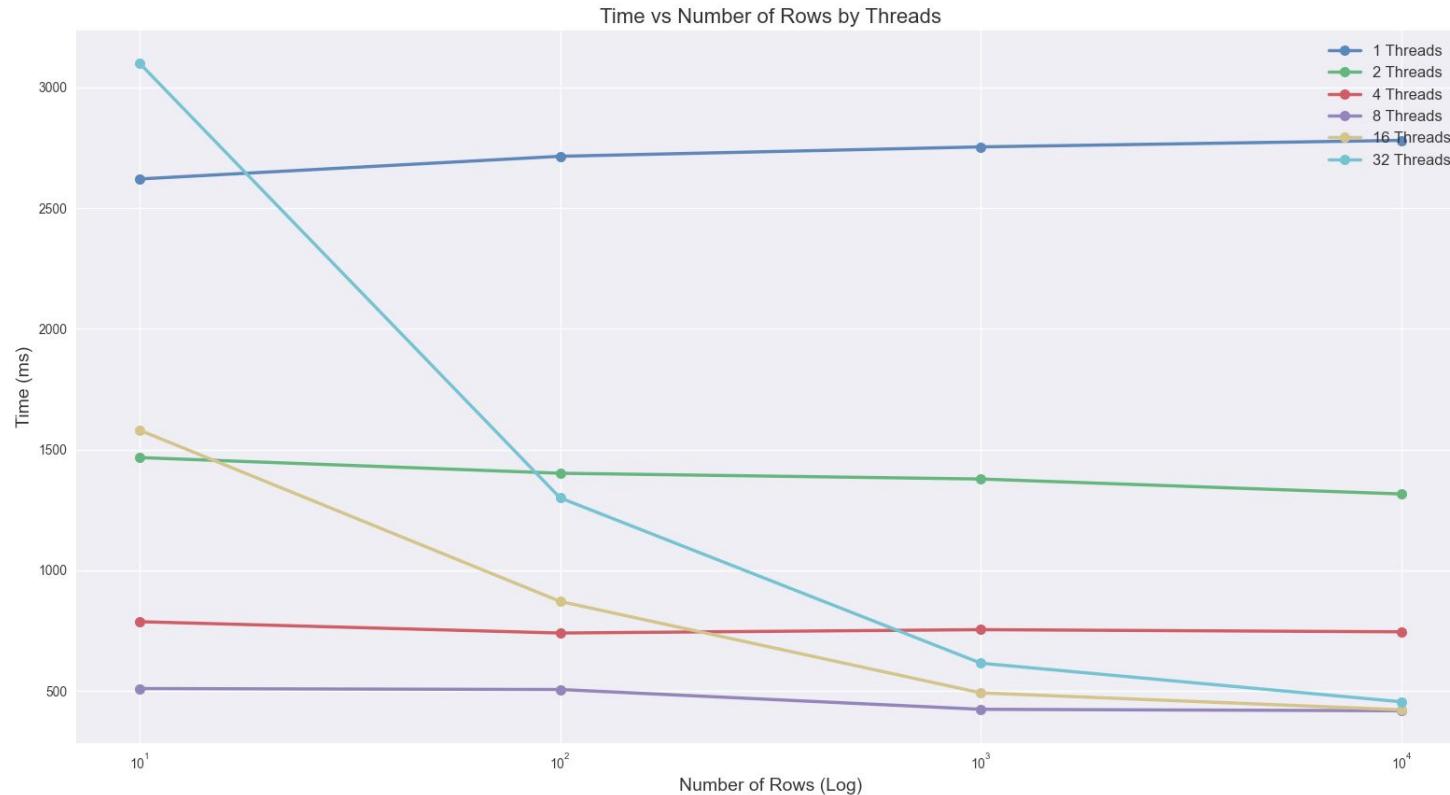
# SIMD

- SIMD used for database searching
- Database size padded to be divisible by 8
- Utilizes AVX2 comparator operator

Loads value to search for	→	12171	12171	12171	12171	12171	12171	12171	12171
Loads 8 search values at a time	→	4613	53130	3230	33131	97100	12171	32304	99200
Generates mask	→	0	0	0	0	0	1	0	0



# Performance Results



**Tested on  
100,000 random  
access read/write  
transactions**



# Performance Results

---

- Databases with low cardinality or high simultaneous ID modifications should remain sequentially accessed
- Repeatable Read with Read/Write Locking best suited for high cardinality databases with lots of random access
- As more threads are added, probability of collision increases which diminishes the benefit → Scale number of threads based on database demands



# Future Work

---

- Develop more efficient locking system
  - Use queue for pending lock requests from concurrent transactions
- Implement logging in database system
  - Ensure atomicity: undo, redo
- Crash recovery
  - Add checkpoint entries to log
  - Use write-ahead logging to restore database after crash

# Questions?

# PROPER AND IMPROPER USE OF A POWERPOINT

---

- You should:

- Use the template
- Use font “Arial”
- Use good-quality images
- Make punctuation and capitalization consistent
- Use bullets
- Include key points and concepts
- Use diagrams to illustrate complex concepts

- You should not:

- Use PowerPoint as a handout, putting all spoken words on slides
- Read your slides or speak to them
- Use whole sentences or paragraphs on your slides (except for quotes)
- Use complicated graphics or charts
- Use flashy transitions and animations
- Use poor-quality or web-napped images



# BOXED CONTENT SLIDE

## Box 1 Header

- Bullet text here...

## Box 2 Header

- Bullet text here...

## Box 3 Header

- Bullet text here...



# BOXED CONTENT SLIDE 2

## Box 1 Header

- Bullet text here...

## Box 2 Header

- Bullet text here...



# SLIDE LAYOUT OPTIONS

- Title slide
  - Used for presentation title only
- Closing slide
  - No content to be added
- Content w/ bullets
  - Style to be used for text (bullets, etc.)
- Photo or chart slides
  - Layout using large image or one to the left or right with the footer
- Divider slide (optional)
  - Use to separate topics or speakers



# IMAGE SLIDES



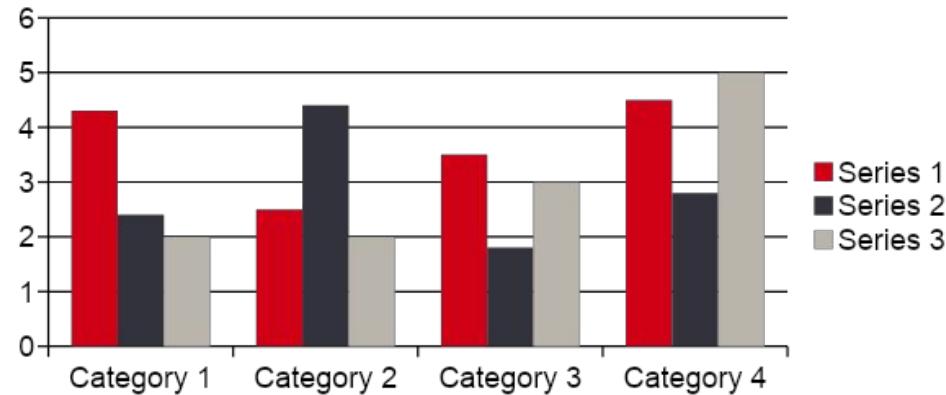
- Small collection of Rensselaer images are placed throughout the deck
  - Backgrounds for section headers
  - Backgrounds for quotes
  - Any high-quality image or chart can be used
  - Do not use as a background for bulleted lists



# CHARTS AND GRAPHS SLIDES

## Including charts and graphs

- In the insert tab, you can create your own charts in PPT
- Multiple options of charts
- Charts and graphs can help the viewer easily understand complex information



# TEXT ON COLOR PICTURES

---

- Level 2 – Arial 18 pt
- Level 3 – Arial 14 pt



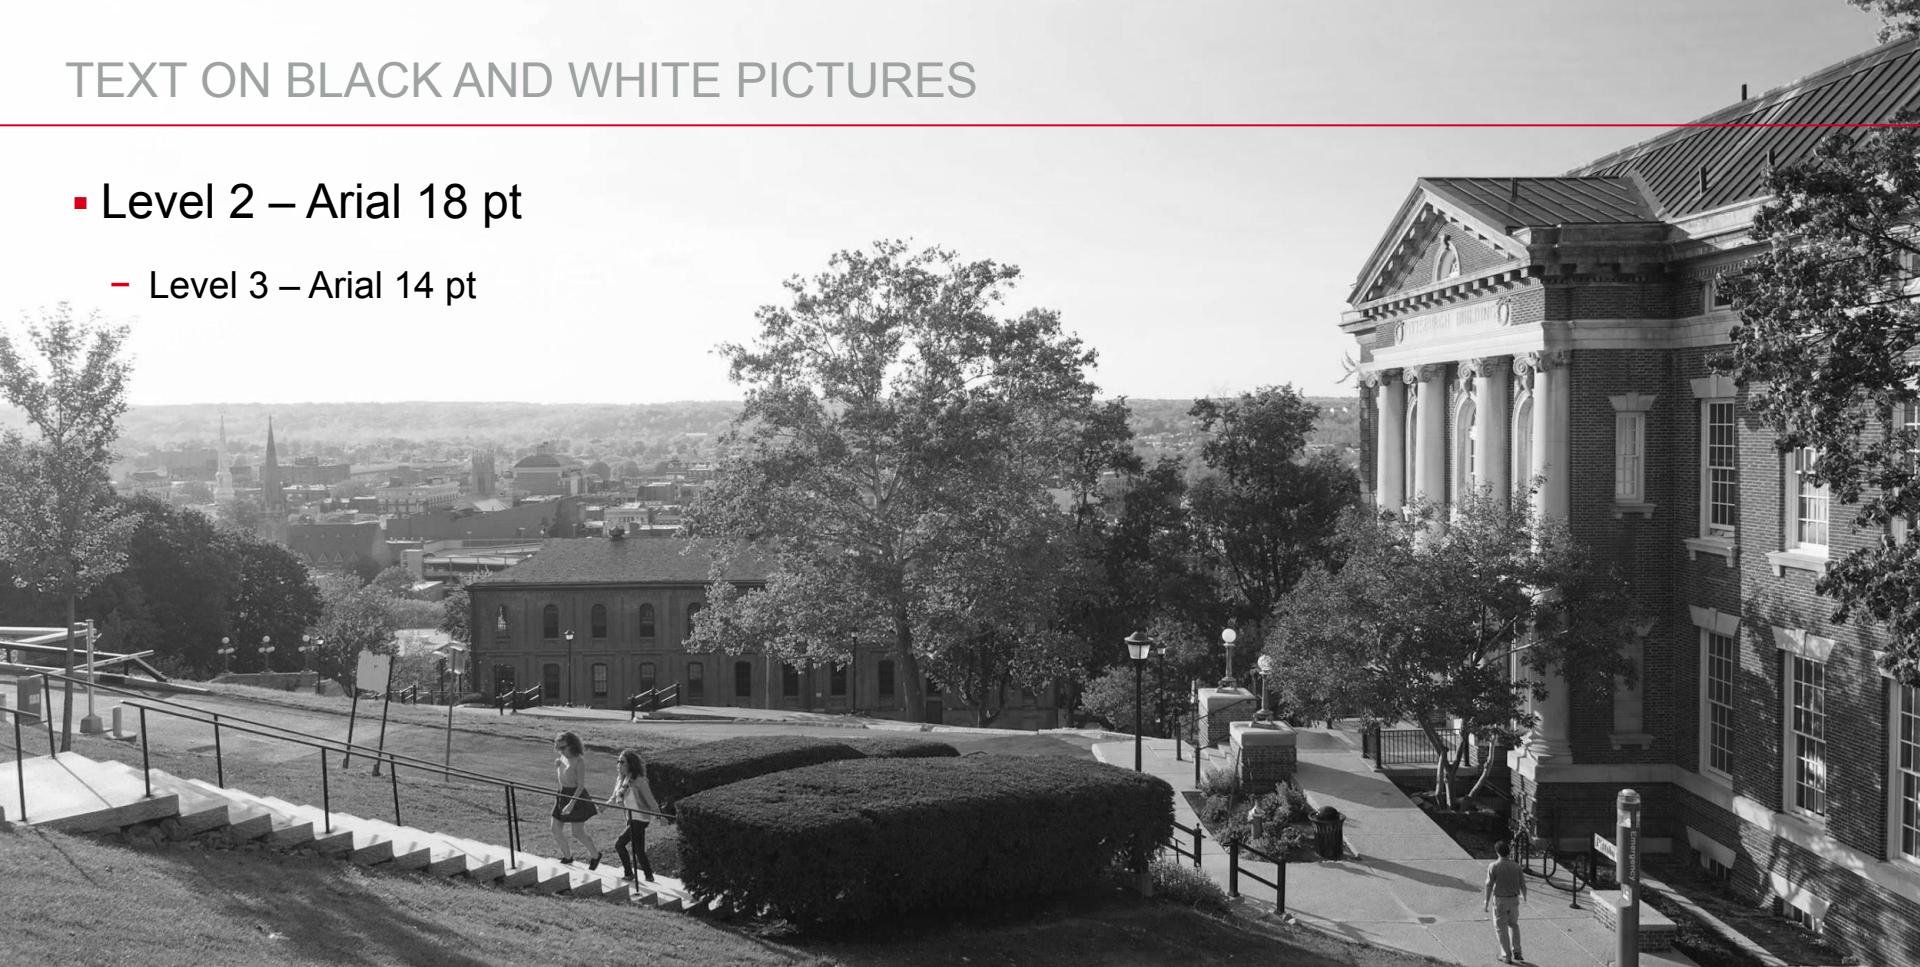
Rensselaer

{ INSERT TITLE HERE }

# TEXT ON BLACK AND WHITE PICTURES

---

- Level 2 – Arial 18 pt
- Level 3 – Arial 14 pt



Rensselaer

{ INSERT TITLE HERE }

# BRAND STANDARDS

---

To get more information on brand style, templates and PowerPoint guidelines reach out to [scer.rpi.edu](http://scer.rpi.edu)

- Level 2 – Arial 18 pt
  - Level 3 – Arial 14 pt





Rensselaer

{ INSERT TITLE HERE }



Rensselaer

{ INSERT TITLE HERE }







Rensselaer

{ INSERT TITLE HERE }

39

11/30/2018



Rensselaer

{ INSERT TITLE HERE }



Rensselaer

{ INSERT TITLE HERE }



Rensselaer

{ INSERT TITLE HERE }



Rensselaer

{ INSERT TITLE HERE }



# Rensselaer

**why not change the world?®**



# Rensselaer

**why not change the world?®**